

Deadlocks

Deadlock problem. Necessary conditions for deadlock.
Resource allocation graphs. Java deadlock examples.
Deadlock prevention. Deadlock avoidance (safe state and Banker's algorithm). Deadlock detection and recovery.

OS5: 22/2/2018


Textbook (SGG): Ch. 7.1-7.4, 7.5.1, 7.5.3, 7.6.2, 7.7

The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set (*circular wait*: simplest process P waiting for Q; Q waiting for P)
- Example 1
 - System has 2 disk drives
 - P_1 and P_2 each hold one disk drive and each needs the other one
- Example 2: (binary) semaphores A and B , initialized to 1

P_0	P_1
acquire(A);	acquire(B);
acquire(B);	acquire(A);

Bridge Crossing Example



- Traffic only in one direction
- Each section of a bridge can be viewed as a resource
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
 - Several cars may have to back up in general
- Starvation is possible (cars in one direction only keep going)
- Note: most operating systems do not prevent or resolve deadlock completely (users will deal with it when needed)

General System Model

- Resource types R_1, R_2, \dots, R_m
 - ▶ Physical interpretation of resource types: CPU, memory, I/O devices, etc
- OS has W_i instances (or units) of each resource type R_i that it can allocate to requesting processes
- Each process utilizes a resource as follows:
 - **request** (request one instance of a resource from OS; e.g., get access to a printer)
 - **use** (use the acquired resource privately to do its work; e.g., print a file to the acquired printer)
 - **release** (return an acquired instance of a resource back to the OS; e.g., give back the acquired printer so the OS can give this printer to another process)

Necessary Conditions for Deadlock

Deadlock *can* occur (but may not necessarily occur) if these four conditions hold *simultaneously*:

- **Mutual exclusion:** Only one process at a time can use a resource
- **Hold and wait:** A process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** A resource can be released only voluntarily by the process holding it, after process has completed its task (hence use of the resource)
- **Circular wait:** There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0

NB: These conditions are *necessary*, but *not sufficient*, for deadlock

Implication: In principle, we can solve deadlock by removing *any* of these conditions

Resource Allocation Graph

A *directed graph*, with set of vertices/nodes V and set of edges E

- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set of all the processes in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set of all the resource types in the system
- **request edge**: directed edge $P_i \rightarrow R_j$
 - P_i wants to acquire an instance of R_j
- **assignment edge**: directed edge $R_j \rightarrow P_i$
 - An instance of R_j is being held by P_i

Resource-Allocation Graph (cont'd)


- Process (first type of nodes/vertices)




- Resource type with 4 instances (second type of nodes)




- P_i requests an instance of R_j (first type of directed edge)



- P_i is holding an instance of R_j (second type of directed edge)



Example Resource Allocation Graph




Is this system deadlocked? *Hint:* Is there circular wait?

Resource Allocation Graph with a Deadlock


P_1 holds an instance of R_2 , needs an instance of R_1 to continue execution so that it may complete its task later

This graph *has* a cycle



Can any of P_1, P_2, P_3 acquire the resource it wants to complete its task?
 Hint: The processes won't release the resources they are holding until they can continue execution and complete their respective tasks.

Graph with A Cycle but No Deadlock



Multiple instances of resource: Cycle is necessary, but not sufficient, condition for deadlock

This graph has a cycle. But what is a possible completion sequence of these processes?

Summary: Graph Cycles and Deadlock

- If graph contains no cycles \Rightarrow no deadlock
- If graph has a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, then *possibility* of deadlock

Java Deadlock Example

```
class A implements Runnable
{
    private Lock first, second;

    public A(Lock first, Lock second) {
        this.first = first;
        this.second = second;
    }

    public void run() {
        try {
            first.lock();
            // do something
            second.lock();
            // do something else
        }
        finally {
            first.unlock();
            second.unlock();
        }
    }
}
```

Thread A

```
class B implements Runnable
{
    private Lock first, second;

    public A(Lock first, Lock second) {
        this.first = first;
        this.second = second;
    }

    public void run() {
        try {
            second.lock();
            // do something
            first.lock();
            // do something else
        }
        finally {
            second.unlock();
            first.unlock();
        }
    }
}
```

Thread B

Note: each runnable object needs to get *two* locks to do its work

Java Deadlock Example

```
public static void main(String arg[]) {  
    Lock lockX = new ReentrantLock();  
    Lock lockY = new ReentrantLock();  
  
    Thread threadA = new Thread(new A(lockX,lockY));  
    Thread threadB = new Thread(new B(lockX,lockY));  
  
    threadA.start();  
    threadB.start();  
}
```

lockX and **lockY** are binary locks, initially available. How can deadlock occur? (Demonstrate an interleaving of the two threads' execution.)

Hint: Each runnable object in previous slide runs as a separate thread;
These threads acquire the locks **lockX** and **lockY** in *different* orders

Methods for Handling Deadlocks

■ Deadlock *avoidance*

- Before granting a resource request (even if request is valid and the requested resources are now available), check that the request will not cause the system to enter a deadlock state (not just no deadlock immediately, but *not even later*)
- Requires advance knowledge of *future* resource needs (e.g., Banker's algorithm in your OS Lab #3)

■ Deadlock *detection and recovery*

- Detect deadlock after the fact, then recover from it (e.g., preempting held resources and rolling back processes)

■ Deadlock *prevention*

- Impose conditions on resource requests to ensure that a valid request can *never* cause the system to enter a deadlock state *by design* (so no need to check the runtime system state as in deadlock avoidance)

■ Real-world OS (e.g., UNIX) may not handle deadlocks completely (i.e., possible that users will see a deadlock; then they'll deal with it)

Deadlock Avoidance

Requests can lead to deadlock (in the future); don't grant those requests (now) even if the requested resources are currently available

- Need additional a priori information: Simplest and most useful model requires each process to declare (in advance) the *maximum number* of resources of each type that it will ever need.
- At time of resource request, avoidance algorithm examines the system's resource allocation state to ensure that granting the request will *never* lead to a *circular-wait* condition later.
- Resource allocation state is defined by the numbers of currently available and allocated resources, and the maximum number of resources that each process may need in the future.


Safe State

- When a process requests an available resource, system must decide if granting the request will leave the system in a safe state.
- System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of all the processes in the system such that for each P_i , the resources that P_i will ever need can be satisfied by the currently available resources *plus* the resources held by all the preceding P_j , with $j < i$.
- Rationale:
 - If P_i 's resource needs are not immediately available, then P_i can wait until all the preceding P_j processes have finished.
 - When the preceding processes all finished, P_i can obtain its needed resources, do its job, *return its allocated resources*, and finish.
 - When P_i finishes, P_{i+1} can obtain its needed resources and finish, and so on.
 - I.e., $\langle P_1, P_2, \dots, P_n \rangle$ defines a feasible order for all the processes to finish.

Basic Facts

- If system is in safe state \Rightarrow no deadlocks
- If system is in unsafe state \Rightarrow *possibility* of deadlock
- Avoidance \Rightarrow ensure that system will never enter an unsafe state.
 - So we'll grant a resource request only if *after* granting the request, the system will still be in a safe state.
 - We need an algorithm to decide if a system is in a safe state (Slide 7.22)

Safe, Unsafe, Deadlock States



Avoidance algorithms

- Single instance of a resource type
 - Use a resource-allocation graph
 - We won't cover this algorithm (limited applicability)
- Multiple instances of a resource type
 - Use the Banker's algorithm (your Lab #3)
 - Subsumes the single-instance problem above

Banker's Algorithm

- Multiple instances of each resource
- Each process P must *a priori* declare its maximum needs (i.e., maximum number of instances of each resource type that P will ever need)
- When a process requests a resource, it may have to wait
- After a process P got all its resources, it must return them within a finite amount of time (when P finishes its task)

Data Structures for Banker's Algorithm

Let n = number of processes, and m = number of resources types

- **Available:** Vector (i.e., 1D array) of length m . If $Available[j] = k$, there are k instances of resource type R_j available
- **Max:** $n \times m$ matrix (i.e., 2D array). If $Max[i,j] = k$, then process P_i may request at most k instances of resource type R_j
- **Allocation:** $n \times m$ matrix. If $Allocation[i,j] = k$ then P_i is currently allocated k instances of R_j
- **Need:** $n \times m$ matrix. If $Need[i,j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:
 $Work = Available$ // NB: *Work*, *Available* are both arrays
 $Finish[i] = false$ for $i = 0, 1, \dots, n - 1$
2. Find an i such that:
 (a) $Finish[i] = false$
 (b) $Need[i] \leq Work$
 If no such i exists, go to Step 4
3. $Work = Work + Allocation[i]$ // $Allocation[i]$ is i -th row of *Allocation* matrix
 $Finish[i] = true$
 go to Step 2
4. If $Finish[i] == true$ for all i , then the system is in a safe state; otherwise it's unsafe

Algorithm for Granting Request by Process P_i

$Request[i]$ = request vector for process P_i . If $Request[i, j] = k$, then process P_i wants k instances of resource type R_j

1. If $Request[i] \leq Need[i]$ go to Step 2. Otherwise, raise error, since process has exceeded its maximum claim.
2. If $Request[i] \leq Available$, go to Step 3. Otherwise P_i must wait, since the resources are not immediately available.
3. Try to allocate the requested resources to P_i by updating the resource allocation state as follows (i.e., assume we grant the new request):

$$Available = Available - Request[i];$$

$$Allocation_i = Allocation[i] + Request[i];$$

$$Need[i] = Need[i] - Request[i];$$

- If *new state is safe* \Rightarrow the resources are allocated to P_i
- If *new state is unsafe* $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored (i.e., new request not granted after all)

Activity 5.1: Banker's Algorithm

- 5 processes P_0 through P_4 ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances)

Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

- Is the system safe? If so, give an execution sequence of the processes that demonstrates the safety. If not, why not?

Activity 5.1 (cont'd)

- Given the original system state in the previous slide, P_1 now requests (1,0,2)
- First, check that $\text{Request} \leq \text{Available}$, i.e., $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$
- Second, now assume we grant this request, which will update the data structures as follows:

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 1	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies the safety condition (i.e., system will remain safe after granting (1,0,2) to P_1)
- Can we now *further* grant (i) request (3,3,0) by P_4 ; or (ii) request (0,2,0) by P_0 ?

Deadlock Detection

- Allow system to enter deadlock state
- Detect occurrence of deadlock by *detection algorithm*
- Recover from the detected deadlock

Deadlock Detection: Multiple Instances

- **Available:** A vector (i.e., 1D array) of length m indicates the number of available resources of each type.
- **Allocation:** An $n \times m$ matrix (i.e., 2D array) defines the number of resources of each type currently allocated to each process.
- **Request:** An $n \times m$ matrix indicates the current request of each process. If $Request[i,j] = k$, then process P_i is requesting k more instances of resource type R_j .

NB: The deadlock *detection* algorithm (next slide) detects if the system is *already* deadlocked. It's similar to the safety algorithm (Slide 7.22) used by the Banker's algorithm for deadlock avoidance, but the safety algorithm detects if the system *can* become deadlocked *in the future*.

Detection Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:
Work = *Available*
For $i = 1, 2, \dots, n$, if *Allocation*[i] $\neq 0$, then
Finish[i] = false; else *Finish*[i] = true
2. Find an index i such that both:
(a) *Finish*[i] == false
(b) *Request*[i] \leq *Work*. // What i requests now is available if all its preceding processes
// finish; i isn't deadlocked already
If no such i exists, go to Step 4
3. *Work* = *Work* + *Allocation*[i]. // Assume i will finish and return the resources it holds now
Finish[i] = true
go to Step 2
4. If *Finish*[i] == false for some i , then the system is (already) in deadlock state.
Moreover, if *Finish*[i] == false, then P_i is (already) deadlocked.

What is the complexity (in terms of m and n) of the algorithm?

Example of Detection Algorithm

- Five processes P_0 through P_4
- Three resource types: A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $Finish[i] = \text{true}$ for all i

Detection Example (cont'd)

- Given the original system state in the previous slide, P_2 now requests an additional instance of type C, so that we have these requests:

	<u>Request</u>		
	A	B	C
P_0	0	0	0
P_1	2	0	2
P_2	0	0	1
P_3	1	0	0
P_4	0	0	2

- What's the new state of the system?
 - Can assume that P_0 will return the resources it holds, but there'll still be insufficient resources to satisfy any other processes' requests
 - Deadlock already exists, consisting of processes P_1 , P_2 , P_3 , and P_4

Deadlock Recovery: Abort Processes

- Abort all deadlocked processes (resources held are preempted)
- Abort one process at a time until the deadlock cycle is eliminated
- Which order should we choose to abort?
 - Priority of the process.
 - How long process has computed, and how much longer until completion?
 - Resources the process has used,
 - Resources the process needs to complete.
 - How many processes will system need to abort?
 - Is process interactive or batch?
 - etc ...

Deadlock Prevention

Constrain the ways requests can be made, in order to disallow “hold-and-wait,” or “no preemption,” or “circular wait” (in prevention, valid requests won’t cause deadlock *by design*, no need to check for safety of runtime system state)

- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none (i.e., previously held resources must be released before new resources are requested)
 - OS must support a new system call for acquiring multiple resources *at the same time*
 - ▶ E.g., for semaphores, define a new “acquireAll([list of semaphores])” system call for acquiring all the semaphores in the argument list, in addition to normal acquire() system call
 - Disadvantages
 - ▶ Low resource utilization if a process has to acquire in the beginning all the resources it’ll ever need – why?
 - ▶ Starvation becomes more likely – Process P needs two resources R1 and R2; R1 is available from time to time, similarly R2; but R1 and R2 are never available *at the same time*

Deadlock Prevention (cont'd)

- **No Preemption** – if a process P (holding some resources already) requests another resource that cannot be immediately granted, then OS will force P to release all the resources it already holds
 - Preempted resources are added to the list of resources for which the process needs to wait for again
 - Process will be restarted only when it can obtain all the resources it needs (both the old ones preempted and the new one requested)
 - Disadvantages: preemption costs and starvation
- **Circular Wait** – impose a total ordering of all the resource types, and require that each process requests resources according to that order
 - Say: give an id for each resource type (e.g., 1 for disk, 2 for printer, etc); process can only request resource types in *increasing* id number (i.e., can request printer while holding disk, but *not* vice versa)
 - Disadvantage: Burden on programmer to ensure the order by design, without unnecessarily sacrificing utilization